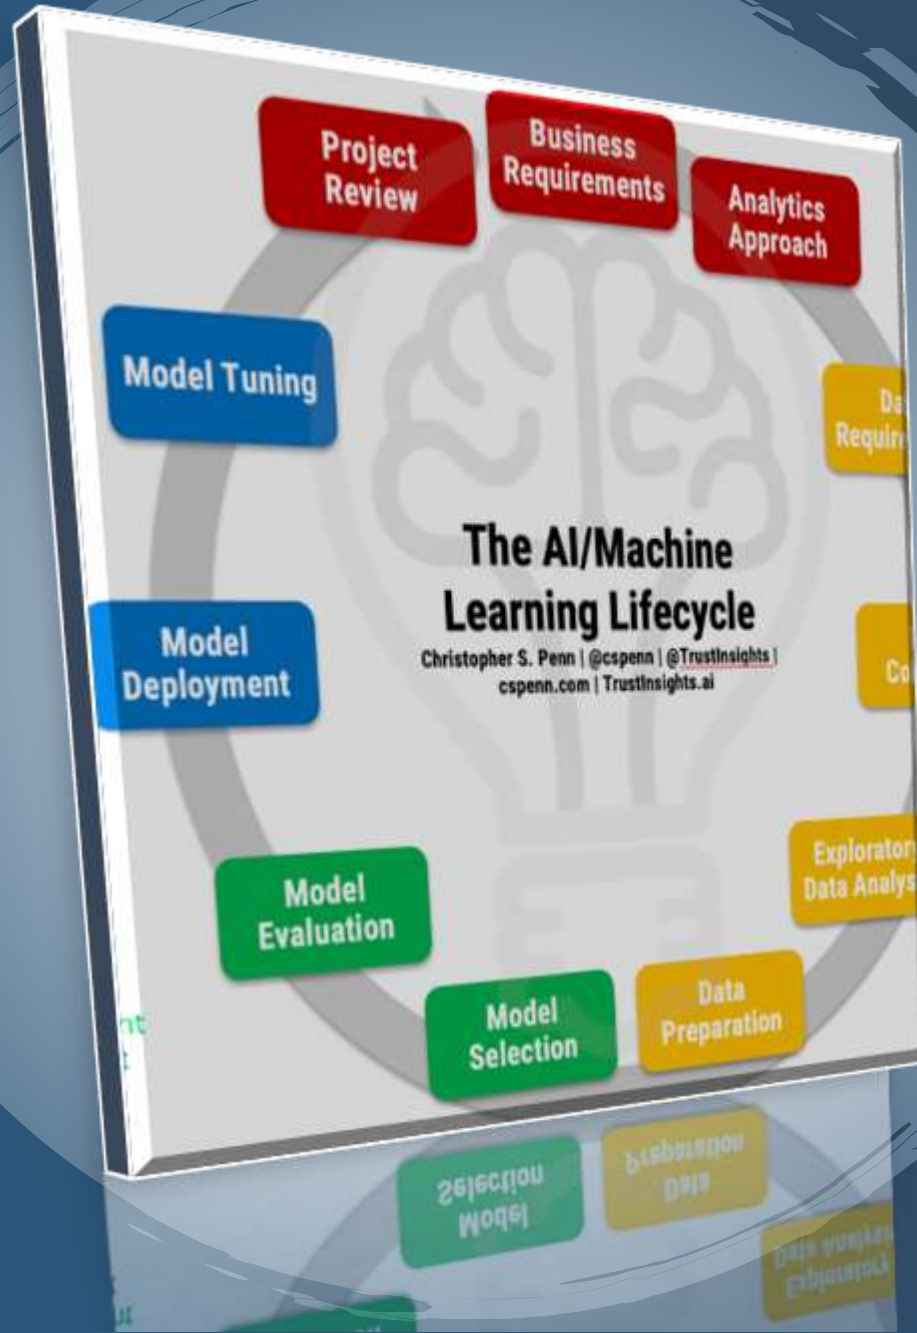


LAPTOP PRICE PREDICTION

FOR SMART TECH CO



WHAT IS MACHINE LEARNING

MACHINE LEARNING IS A BRANCH OF ARTIFICIAL INTELLIGENCE THAT DEALS WITH IMPLEMENTING APPLICATIONS THAT CAN MAKE A FUTURE PREDICTION BASED ON PAST DATA.

THERE ARE VARIOUS STEPS INVOLVED IN BUILDING A MACHINE LEARNING PROJECT BUT NOT ALL THE STEPS ARE MANDATORY TO USE IN A SINGLE PROJECT, AND IT ALL DEPENDS ON THE DATA.



• WHY WE NEED LAPTOP PREDICTION MODEL ?

- Consumers often face the challenge of navigating through various options to find a laptop that meets their requirements and budget constraints. Additionally, market fluctuations and the rapid pace of technological advancements contribute to the complexity of understanding and predicting laptop prices.



LAPTOP PREDICTION MODEL

Project Overview:

Smart Tech Co. has partnered with our data science team to develop a robust machine learning model that predicts laptop prices accurately. As the market for laptops continues to expand with a myriad of brands and specifications, having a precise pricing model becomes crucial for both consumers and manufacturers



Table of Contents

- 1.Describing Problem Statement
- 2.Overview about dataset
- 3.Data Cleaning
- 4.Feature engineering
- 5.Exploratory Data Analysis
- 6.Model training, Model testing
- 7.Machine learning Modeling
- 8.Questions related to our observation
- 9.ML web app development
- 10.Deployment Machine learning app

Problem Statement

The problem statement is that if any user wants to buy a laptop then our application should be compatible to provide a tentative price of laptop according to the user configurations.



SOLUTION

WE WILL MAKE A PROJECT FOR LAPTOP

PRICE PREDICTION.



BASIC UNDERSTANDING OF LAPTOP PRICE PREDICTION DATA

LAPTOP PRICE PRITION MODEL

```
# importing Libraries
```

```
import numpy as np
import pandas as pd
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
##Importing files
```

```
df = pd.read_csv('laptop.csv')
```

```
# Understanding the dataset.
```

```
df.head()
```

	Unnamed: 0.1	Unnamed: 0	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price
0	0	0.0	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8GB	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37kg	71378.6832
1	1	1.0	Apple	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8GB	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34kg	47895.5232

[**df.shape**] # Checking no of rows and no columns.

In this dataset we have; ### ROWS = 1308 ### COLUMNS = 13

Descriptive statistical analysis.
df.describe(include='all')

Descriptive statistical analysis
df.describe(include='all')

	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price
count	1273	1273	1273	1273	1273	1273	1273	1273	1273	1273	1273.000000
unique	19	6	25	40	118	10	40	110	9	189	NaN
top	Lenovo	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8GB	256GB SSD	Intel HD Graphics 620	Windows 10	2.2kg	NaN
freq	290	710	640	495	183	601	401	271	1047	111	NaN
mean	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	59955.814073
std	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	17332.251005
min	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	8270.720000
25%	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	31814.720000
50%	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	52161.120000
75%	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	78333.387200
max	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	324954.720000

df.info()

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 13 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   Unnamed: 0.1        1303 non-null  int64  
 1   Unnamed: 0          1273 non-null  float64
 2   Company             1273 non-null  object  
 3   TypeName            1273 non-null  object  
 4   Inches              1273 non-null  object  
 5   ScreenResolution    1273 non-null  object  
 6   Cpu                 1273 non-null  object  
 7   Ram                 1273 non-null  object  
 8   Memory              1273 non-null  object  
 9   Gpu                 1273 non-null  object  
10  OpSys               1273 non-null  object  
11  Weight              1273 non-null  object  
12  Price               1273 non-null  float64
dtypes: float64(2), int64(1), object(10)
memory usage: 132.5+ KB
```

Checking no of rows and no columns.
[df.shape]

In this dataset we have;
ROWS = 1308
COLUMNS = 13

remove nan values
df.dropna(inplace=True)

```
[177]: # remove nan values
df.dropna(inplace=True)

[178]: df.isnull().sum()

[178]: Unnamed: 0.1      0
      Unnamed: 0      0
      Company         0
      TypeName        0
      Inches          0
      ScreenResolution 0
      Cpu             0
      Ram             0
      Memory          0
      Gpu             0
      OpSys           0
      Weight          0
      Price           0
      dtype: int64
```


Data cleaning & Feature Engineering

```
## Drop unwanted columns from the dataframe  
df.drop(columns=['Unnamed: 0.1', 'Unnamed: 0'], axis=1, inplace=True)
```

```
## Now, check the data frame again df.head()
```

```
## Drop unwanted columns from the dataframe  
df.drop(columns=['Unnamed: 0.1', 'Unnamed: 0'], axis=1, inplace=True)  
  
## Now, check the dataframe again  
df.head()
```

	Company	Type Name	Inches	Screen Resolution	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price
0	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8GB	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37kg	71378.6832
1	Apple	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8GB	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34kg	47895.5232
2	HP	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8GB	256GB SSD	Intel HD Graphics 620	No OS	1.86kg	30636.0000
3	Apple	Ultrabook	15.4	IPS Panel Retina Display 2880x1800	Intel Core i7 2.7GHz	16GB	512GB SSD	AMD Radeon Pro 455	macOS	1.83kg	135195.3360
4	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 3.1GHz	8GB	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37kg	96095.8080

Now, you can see that ram is object, we will change it to [int].

Weight is also object, I will change it to [float].

Change price dtype into [float].

```
# Replace the 'empty' string and '?' with NaN and then drop rows with NaN in 'Ram' and 'Weight' columns
df.replace(['', '?'], np.nan, inplace=True)
df.dropna(subset=['Ram', 'Weight'], inplace=True)

# Now convert to integer and float
df['Ram'] = df['Ram'].str.replace('GB', '').astype('int32')
df['Weight'] = df['Weight'].str.replace('kg', '').astype('float32')

# Remove non-numeric characters from 'Inches' column
df['Inches'] = df['Inches'].replace('?', float('nan'))

# convert 'Inches' column to float
df['Inches'] = df['Inches'].astype(float)

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1272 entries, 0 to 1302
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Company         1272 non-null   object
1   TypeName        1272 non-null   object
2   Inches          1271 non-null   float64
3   ScreenResolution 1272 non-null   object
4   Cpu             1272 non-null   object
5   Ram             1272 non-null   int32
```

```
183: print(df['Ram'].dtype)
      object
184: print(df['Weight'].dtype)
      object
185: print(df['Price'].dtype)
      float64
```

`df['Company'].value_counts()`

```
[191]: df['Company'].value_counts()
```

```
[191]: Company
Lenovo      298
Dell        286
HP          266
Asus        156
Acer        103
MSI         53
Toshiba     47
Apple       21
Samsung     9
Razer       7
Mediacom    7
Microsoft   6
Xiaomi      4
Vero        4
Chuvi       3
Google      3
LG          3
Huawei       2
Fujitsu     2
Name: count, dtype: int64
```

`df['TypeName'].value_counts()`

```
df['TypeName'].value_counts()
```

```
TypeName
Notebook      710
Gaming        203
Ultrabook     190
2 in 1 Convertible 116
Workstation   29
Netbook       24
Name: count, dtype: int64
```

```
df['ScreenResolution'].value_counts()
```

```
ScreenResolution
Full HD 1920x1080      494
1366x768               274
IPS Panel Full HD 1920x1080 226
IPS Panel Full HD / Touchscreen 1920x1080 52
Full HD / Touchscreen 1920x1080 45
1600x900               23
Touchscreen 1366x768    16
Quad HD+ / Touchscreen 3200x1800 14
IPS Panel 4K Ultra HD 3840x2160 12
IPS Panel 4K Ultra HD / Touchscreen 3840x2160 11
4K Ultra HD / Touchscreen 3840x2160 9
4K Ultra HD 3840x2160    7
IPS Panel 1366x768      7
IPS Panel Retina Display 2560x1600 6
IPS Panel Quad HD+ / Touchscreen 3200x1800 6
Touchscreen 2560x1440    6
IPS Panel Retina Display 3200x1440 6
```

`df['Screen Resolution'].value_counts()`

Checking data type after dtype modification.

```
## Checking data type after dtype modification.
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1272 entries, 0 to 1302
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Company                1272 non-null   object
1   TypeName                1272 non-null   object
2   Inches                 1271 non-null   float64
3   ScreenResolution       1272 non-null   object
4   Cpu                    1272 non-null   object
5   Ram                    1272 non-null   int32
6   Memory                 1271 non-null   object
7   Gpu                    1272 non-null   object
8   OpSys                  1272 non-null   object
9   Weight                 1272 non-null   float32
10  Price                  1272 non-null   float64
dtypes: float32(1), float64(2), int32(1), object(7)
memory usage: 109.3+ KB
```

Feature Engineering over screen resolution column

There are some hidden specifications in Screen resolution column which we have to find out and create a separate column for it.



IPS PANEL RETINA
DISPLAY

`df['IPS'] = df['Screen Resolution'].apply(lambda x: 1 if
'IPS' in x else 0)`



TOUCH SCREEN

`df['Screen Resolution'].apply(lambda x:1 if
'Touchscreen' in x else 0)`

df.sample(10)

```
[200]: df.sample(10)
```

[200]:	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	IPS	Touchscreen
635	Asus	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7300HQ 2.5GHz	8	1TB HDD	Nvidia GeForce GTX 1050	Windows 10	1.99	48304.7136	0	0
1057	Acer	Notebook	15.6	1366x768	AMD A8-Series 7410 2.2GHz	8	1TB HDD	AMD Radeon R5	Windows 10	2.40	23922.7200	0	0
1260	Lenovo	2 in 1 Convertible	14.0	Full HD / Touchscreen 1920x1080	Intel Core i5 6200U 2.3GHz	4	128GB SSD	Intel HD Graphics 520	Windows 10	1.80	44382.7728	0	1
1297	Asus	Notebook	15.6	1366x768	Intel Core i7 6500U 2.5GHz	4	500GB HDD	Nvidia GeForce 920M	Windows 10	2.20	38378.6496	0	0
1109	Asus	Gaming	15.6	IPS Panel Full HD 1920x1080	Intel Core i7 6700HQ 2.6GHz	16	128GB SSD + 1TB HDD	Nvidia GeForce GTX 960M	Windows 10	2.59	71341.9200	1	0
1249	Dell	2 in 1 Convertible	13.3	Quad HD+ / Touchscreen 3200x1800	Intel Core i5 7Y54 1.2GHz	8	256GB SSD	Intel HD Graphics 615	Windows 10	1.24	96596.6400	0	1
164	Acer	Notebook	15.6	1366x768	Intel Celeron Dual Core N3350 1.1GHz	4	1TB HDD	Intel HD Graphics 500	Windows 10	2.10	18541.4400	0	0
578	MSI	Gaming	17.3	Full HD 1920x1080	Intel Core i7 7820HK 2.9GHz	16	512GB SSD + 1TB HDD	Nvidia GeForce GTX 1070	Windows 10	4.14	145401.1200	0	0

Create two new columns :

```
df['Screen Resolution'].str.split('x',n=1,expand=True)  
##Creating two new columns..x_res,y_res
```

```
df['X_res']=new[0]  
df['Y_res']=new[1]
```

**Screen
Resolution**

We split it into X_res and Y_res

X_res

X resolution

Y_res

Y resolution

df.head()

```
df['X_res'] = new[0]
df['Y_res'] = new[1]
```

```
df.head()
```

	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	IPS	Touchscreen	X_res	Y_res
0	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	1	0	IPS Panel Retina Display 2560	1600
1	Apple	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34	47895.5232	0	0	1440	900
2	HP	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.86	30636.0000	0	0	Full HD 1920	1080
3	Apple	Ultrabook	15.4	IPS Panel Retina Display 2880x1800	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.83	135195.3360	1	0	IPS Panel Retina Display 2880	1800
4	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	1	0	IPS Panel Retina Display 2560	1600

we have to change X_res.## with the help of regular expression we will extract required info from X_res.

```
df['X_res'].str.replace(',','').str.findall(r'(\d+\.?\d+)').apply(lambda x:x[0])
```

we have to change X_res.

with the help of regular expression we will extract required info from X_res.

```
df['X_res'].str.replace(',','').str.findall(r'(\d+\.?\d+)').apply(lambda x:x[0])
```

```
0      2560
1      1440
2      1920
3      2880
4      2560
...
```

```
1298    1920
1299    3200
1300    1366
1301    1366
1302    1366
```

Name: X_res, Length: 1272, dtype: object

We successfully extracted the values.

```
df['X_res'] = df['X_res'].str.replace(',','').str.findall(r'(\d+\.?\d+)').apply(lambda x:x[0])
```

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>

Index: 1272 entries, 0 to 1302

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	Company	1272 non-null	object
1	TypeName	1272 non-null	object
2	Inches	1271 non-null	float64
3	ScreenResolution	1272 non-null	object
4	Cpu	1272 non-null	object
5	Ram	1272 non-null	int32
6	Memory	1271 non-null	object
7	Gpu	1272 non-null	object
8	OpSys	1272 non-null	object
9	Weight	1272 non-null	float32
10	Price	1272 non-null	float64
11	IPS	1272 non-null	int64
12	Touchscreen	1272 non-null	int64
13	X_res	1272 non-null	object
14	Y_res	1272 non-null	object

dtypes: float32(1), float64(2), int32(1), int64(2), object(9)

memory usage: 149.1+ KB

Converting X_res and Y_res into int dtype.

```
## Converting X_res and Y_res into int dtype.
```

```
df['X_res'] = df['X_res'].astype('int')
```

```
df['Y_res'] = df['Y_res'].astype('int')
```

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>

Index: 1272 entries, 0 to 1302

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	Company	1272 non-null	object
1	TypeName	1272 non-null	object
2	Inches	1271 non-null	float64
3	ScreenResolution	1272 non-null	object
4	Cpu	1272 non-null	object
5	Ram	1272 non-null	int32
6	Memory	1271 non-null	object
7	Gpu	1272 non-null	object
8	OpSys	1272 non-null	object
9	Weight	1272 non-null	float32
10	Price	1272 non-null	float64
11	IPS	1272 non-null	int64
12	Touchscreen	1272 non-null	int64
13	X_res	1272 non-null	int32
14	Y_res	1272 non-null	int32

dtypes: float32(1), float64(2), int32(3), int64(2), object

Now we will create a new column i.e, PPI

What is pixel per inch (PPI)?# The term Pixels Per Inch (PPI) commonly refers to the measurement of **pixel density in display screens, including those of computers, laptops, TVs, and smartphones.** # This metric helps you determine the sharpness and clarity of the image you see on the screens of these devices. # PPI is a metric used for all kinds of screens

```
df['ppi'] = (((df['X_res']**2) + (df['Y_res']**2))**0.5/df['Inches']).astype('float')
```

We check correlation to understand more about numeric columns and which are least important.

```
# Select only numeric columns for correlation calculation
numeric_df = df.select_dtypes(include=['number'])

# Calculate correlations with 'Price'
correlations = numeric_df.corr()['Price']

print(correlations)
```

Inches	0.045042
Ram	0.685737
Weight	0.175928
Price	1.000000
IPS	0.255140
Touchscreen	0.189172
X_res	0.557584
Y_res	0.554104
ppi	0.469329

Name: Price, dtype: float64



Now we will remove some columns which are of no use in future.
we will drop Inches,ScreenResolution,X_res,Y_res.

```
df.drop(columns=['ScreenResolution'],inplace=True)
```

```
df.drop(columns=['X_res','Y_res'],inplace=True)  
df.drop(columns=['Inches'],inplace=True)
```

```
df.head()
```

	Company	TypeName	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	IPS	Touchscreen	ppi
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	1	0	226.983005
1	Apple	Ultrabook	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34	47895.5232	0	0	127.677940
2	HP	Notebook	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.86	30636.0000	0	0	141.211998
3	Apple	Ultrabook	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.83	135195.3360	1	0	220.534624
4	Apple	Ultrabook	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	1	0	226.983005

Now we will focus on cpu and gpu

cpu

```
df['Cpu'].value_counts()
```

```
df['Cpu'].apply(lambda x:x.split()[0:3])
```

```
## I will create a new column [Cpu Name]
```

```
df['Cpu Name'] = df['Cpu'].apply(lambda x:" ".join(x.split()[0:3]))
```

```
df['Cpu'].value_counts()
```

```
Cpu
Intel Core i5 7200U 2.5GHz    183
Intel Core i7 7700HQ 2.8GHz   142
Intel Core i7 7500U 2.7GHz    128
Intel Core i7 8550U 1.8GHz     71
Intel Core i5 8250U 1.6GHz     68
...
AMD A9-Series 9420 2.9GHz      1
Intel Core i7 2.2GHz            1
AMD A6-Series 7310 2GHz         1
Intel Atom Z8350 1.92GHz        1
AMD E-Series 9000e 1.5GHz        1
Name: count, Length: 118, dtype: int64
```

```
df['Cpu'].apply(lambda x:x.split()[0:3])
```

```
0      [Intel, Core, i5]
1      [Intel, Core, i5]
2      [Intel, Core, i5]
3      [Intel, Core, i7]
4      [Intel, Core, i5]
...
1298   [Intel, Core, i7]
1299   [Intel, Core, i7]
1300   [Intel, Celeron, Dual]
1301   [Intel, Core, i7]
1302   [Intel, Celeron, Dual]
Name: Cpu, Length: 1272, dtype: object
```

```
df.head()
```

```
[228]:
```

	Company	TypeName	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	IPS	Touchscreen	ppi	Cpu Name
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	1	0	226.983005	Intel Core i5
1	Apple	Ultrabook	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34	47895.5232	0	0	127.677940	Intel Core i5
2	HP	Notebook	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.86	30636.0000	0	0	141.211998	Intel Core i5
3	Apple	Ultrabook	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.83	135195.3360	1	0	220.534624	Intel Core i7
4	Apple	Ultrabook	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	1	0	226.983005	Intel Core i5

Gpu

```
df['Gpu'].value_counts()
```

```
df['Gpu brand'] = df['Gpu'].apply(lambda x:x.split()[0])
```

```
df.head()
```

```
df['Gpu brand'].value_counts()
```

```
df.drop(columns=['Gpu'],inplace=True)
```

```
df.head()
```

```
df['Gpu'].value_counts()
```

```
Gpu
Intel HD Graphics 620    271
Intel HD Graphics 520    181
Intel UHD Graphics 620    65
Nvidia GeForce GTX 1050    64
Nvidia GeForce GTX 1060    48
...
AMD Radeon R5 520         1
AMD Radeon R7             1
Intel HD Graphics 540     1
AMD Radeon 540            1
ARM Mali T860 MP4         1
Name: count, Length: 110, dtype: int64
```

```
df['Gpu brand'] = df['Gpu'].apply(lambda x:x.split()[0])
```

```
df.head()
```

	Company	TypeName	Cpu	Ram	Gpu	OpSys	Weight	Price	IPS	Touchscreen	ppi	Cpu Name	Cpu brand	SSD	HDD	Gpu brand
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	1	0	226.983005	Intel Core i5	Intel Core i5	128	0	Intel
1	Apple	Ultrabook	Intel Core i5 1.8GHz	8	Intel HD Graphics 6000	macOS	1.34	47895.5232	0	0	127.677940	Intel Core i5	Intel Core i5	0	0	Intel
2	HP	Notebook	Intel Core i5	8	Intel HD	No OS	1.86	30636.0000	0	0	141.211998	Intel	Intel	256	0	Intel

```
df['Gpu brand'].value_counts()
```

```
Gpu brand
Intel    702
Nvidia   393
AMD      176
ARM        1
Name: count, dtype: int64
```

```
df.drop(columns=['Gpu'],inplace=True)
```

```
df.head()
```

	Company	TypeName	Cpu	Ram	OpSys	Weight	Price	IPS	Touchscreen	ppi	Cpu Name	Cpu brand	SSD	HDD	Gpu brand
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	macOS	1.37	71378.6832	1	0	226.983005	Intel Core i5	Intel Core i5	128	0	Intel
1	Apple	Ultrabook	Intel Core i5 1.8GHz	8	macOS	1.34	47895.5232	0	0	127.677940	Intel Core i5	Intel Core i5	0	0	Intel
2	HP	Notebook	Intel Core i5 7200U 2.5GHz	8	No OS	1.86	30636.0000	0	0	141.211998	Intel Core i5	Intel Core i5	256	0	Intel
3	Apple	Ultrabook	Intel Core i7 2.7GHz	16	macOS	1.83	135195.3360	1	0	220.534624	Intel Core i7	Intel Core i7	512	0	AMD

Now we will focus on memory

```
df['Memory'].value_counts()
```

```
import re
```

```
# Extract SSD and HDD values using regular expressions
```

```
df['SSD'] = df['Memory'].apply(lambda x: int(re.findall(r'(\d+)GB SSD', str(x))[0]) if re.findall(r'(\d+)GB SSD', str(x)) else 0)
```

```
df['HDD'] = df['Memory'].apply(lambda x: int(re.findall(r'(\d+)GB HDD', str(x))[0]) if re.findall(r'(\d+)GB HDD', str(x)) else 0)
```

df.head()

```
df['Memory'].value_counts()
```

Memory	
256GB SSD	400
1TB HDD	217
500GB HDD	130
512GB SSD	116
128GB SSD + 1TB HDD	92
128GB SSD	74
256GB SSD + 1TB HDD	71
32GB Flash Storage	37
2TB HDD	16
64GB Flash Storage	14
512GB SSD + 1TB HDD	14
1TB SSD	13
256GB SSD + 2TB HDD	10
1.0TB Hybrid	9
256GB Flash Storage	8
16GB Flash Storage	7
32GB SSD	6
128GB Flash Storage	4
180GB SSD	4
512GB SSD + 2TB HDD	3
16GB SSD	3
512GB Flash Storage	2
1TB SSD + 1TB HDD	2
256GB SSD + 500GB HDD	2
128GB SSD + 2TB HDD	2
256GB SSD + 256GB SSD	2
512GB SSD + 256GB SSD	1

```

# Extract SSD and HDD values using regular expressions
df['SSD'] = df['Memory'].apply(lambda x: int(re.findall(r'(\d+)GB SSD', str(x))[0]) if re.findall(r'(\d+)GB SSD', str(x)) else 0)
df['HDD'] = df['Memory'].apply(lambda x: int(re.findall(r'(\d+)GB HDD', str(x))[0]) if re.findall(r'(\d+)GB HDD', str(x)) else 0)

df.head()

```

	Company	TypeName	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	IPS	Touchscreen	ppi	Cpu Name	Cpu brand	SSD	HDD
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	1	0	226.983005	Intel Core i5	Intel Core i5	128	0

df.drop('Memory', axis=1, inplace=True)

	Company	TypeName	Cpu	Ram	Gpu	OpSys	Weight	Price	IPS	Touchscreen	ppi	Cpu Name	Cpu brand	SSD	HDD
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	1	0	226.983005	Intel Core i5	Intel Core i5	128	0
1	Apple	Ultrabook	Intel Core i5 1.8GHz	8	Intel HD Graphics 6000	macOS	1.34	47895.5232	0	0	127.677940	Intel Core i5	Intel Core i5	0	0
2	HP	Notebook	Intel Core i5 7200U 2.5GHz	8	Intel HD Graphics 620	No OS	1.86	30636.0000	0	0	141.211998	Intel Core i5	Intel Core i5	256	0
3	Apple	Ultrabook	Intel Core i7 2.7GHz	16	AMD Radeon Pro 455	macOS	1.83	135195.3360	1	0	220.534624	Intel Core i7	Intel Core i7	512	0
4	Apple	Ultrabook	Intel Core i5 3.1GHz	8	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	1	0	226.983005	Intel Core i5	Intel Core i5	256	0

Operating system

```
df['OpSys'].value_counts()
```

```
OpSys
Windows 10      1046
No OS           63
Linux           61
Windows 7       45
Chrome OS       27
macOS           13
Mac OS X        8
Windows 10 S    8
Android         1
Name: count, dtype: int64
```

```
## we will club it
```

```
def cat_os(inp):
    if inp == 'Windows 10' or inp == 'Windows 7' or inp == 'Windows 10 S':
        return 'Windows'
    elif inp == 'macOS' or inp == 'Mac OS X':
        return 'Mac'
    else:
        return 'Others/No OS/Linux'
```

```
df['os'] = df['OpSys'].apply(cat_os)
```

```
df.head()
```



```
df.head()
```

	Company	TypeName	Cpu	Ram	OpSys	Weight	Price	IPS	Touchscreen	ppi	Cpu Name	Cpu brand	SSD	HDD	Gpu brand	os
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	macOS	1.37	71378.6832	1	0	226.983005	Intel Core i5	Intel Core i5	128	0	Intel	Mac
1	Apple	Ultrabook	Intel Core i5 1.8GHz	8	macOS	1.34	47895.5232	0	0	127.677940	Intel Core i5	Intel Core i5	0	0	Intel	Mac
2	HP	Notebook	Intel Core i5 7200U 2.5GHz	8	No OS	1.86	30636.0000	0	0	141.211998	Intel Core i5	Intel Core i5	256	0	Intel	Others/No OS/Linux
3	Apple	Ultrabook	Intel Core i7 2.7GHz	16	macOS	1.83	135195.3360	1	0	220.534624	Intel Core i7	Intel Core i7	512	0	AMD	Mac
4	Apple	Ultrabook	Intel Core i5 3.1GHz	8	macOS	1.37	96095.8080	1	0	226.983005	Intel Core i5	Intel Core i5	256	0	Intel	Mac

```
df.drop(columns=['OpSys'], inplace=True)
```

EDA of Laptop Price Prediction Dataset

Exploratory analysis is a process to explore and understand the data and data relationship in a complete depth so that it makes feature engineering and machine learning modeling steps smooth and streamlined for prediction.



EDA involves Univariate, Bivariate, or Multivariate analysis.

EDA helps to prove our assumptions true or false. In other words, it helps to perform hypothesis testing.

import the necessary libraries


```
import seaborn as sns  
import matplotlib.pyplot as plt
```

THE EPICENTRE OF EDA IS PRICE


Treating outliers

```
# Filling the null values with median of each  
variable  
df['Ram']=df['Ram'].fillna(df['Ram'].median())  
df['Weight']=df['Weight'].fillna(df['Weight'].media))  
df['Price']=df['Price'].fillna(df['Price'].median())
```

```
sns.boxplot(df['Price'])
```




```
q3=df.describe()['Price']['75%']  
q1=df.describe()['Price']['25%']  
iqr=q3-q1  
upper=q3+1.5*iqr  
lower=q1-1.5*iqr  
df['Price']=df['Price'].clip(upper,lower)
```




Treating outliers

```
import seaborn as sns
sns.boxplot(df['ppi'])
```



```
q3=df.describe()['ppi']['75%']
q1=df.describe()['ppi']['25%']
iqr=q3-q1
upper=q3+1.5*iqr
lower=q1-1.5*iqr
df['ppi']=df['ppi'].clip(upper,lower)
```


```
sns.boxplot(df['ppi'])
```




Treating outliers

```
q3=df.describe()['Ram']['75%']  
q1=df.describe()['Ram']['25%']  
iqr=q3-q1upper=q3+1.5*iqr  
lower=q1-1.5*iqrdf['Ram']=df['Ram'].clip(upper,lower)
```

sns.boxplot(df['Ram'])




sns.boxplot(df['Ram'])






PRICE ANALYSIS

```
sns.displot(df['Price'])
```





```
sns.barplot(x=df['Company'],y=df['Price'])plt.xticks(rotation='vertical')plt.show
```





PRICE ANALYSIS

```
sns.barplot(x=df['TypeName'],y=df['Price'])plt.xticks(rotation='vertical')  
plt.show
```




```
sns.barplot(x=df['Touchscreen'],y=df['Price'])
```




PRICE ANALYSIS

```
sns.barplot(x=df['IPS'],y=df['Price'])
```




```
sns.barplot(x=df['os'],y=df['Price'])plt.xticks(rotation='vertical')plt.show()
```





PRICE ANALYSIS

```
sns.barplot(x=df['Gpubrand'],y=df['Price'])plt.xticks(rotation='vertical')  
plt.show
```



```
sns.barplot(x=df['Ram'],y=df['Price'])plt.xticks(rotation='vertical')  
plt.show
```



MACHINE LEARNING : Model Training & Model Testing

Machine Learning Problem

It is a Regression problem, for a given columns we need to predict the price of laptop.

Performance Metric

1). R² Score


2). Mean Absolute Error

MODEL TRAINING

Normalisation

extract x and y by applying log transformation.
It will increase the value of R2.

```
sns.distplot(np.log(df['Price']))
```



```
from sklearn.model_selection import  
train_test_splitX_train,X_test,y_train,y_test =  
train_test_split(X,y,test_size=0.15,random_state=2)
```

X_train

```
[163]: from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.15,random_state=2)
```

```
[164]: X_train
```

```
[164]:
```

	Company	TypeName	Ram	Weight	IPS	Touchscreen	ppi	Cpu brand	SSD	HDD	Gpu brand	os
735	Lenovo	Notebook	4	1.85	0	0	141.211998	Intel Core i7	0	0	Intel	Windows
22	HP	Notebook	4	1.86	0	0	100.454670	AMD Processor	0	500	AMD	Others/No OS/Linux
811	MSI	Gaming	14	2.90	0	0	127.335675	Intel Core i7	512	0	Nvidia	Windows
283	Lenovo	Notebook	6	2.20	0	0	141.211998	Intel Core i5	256	0	Intel	Windows
765	Acer	Notebook	4	1.60	0	0	117.826530	Intel Core i5	128	0	Intel	Windows
...
479	Toshiba	Notebook	8	1.05	1	0	165.632118	Intel Core i5	256	0	Intel	Windows
309	HP	Notebook	4	1.86	0	0	141.211998	Intel Core i3	0	0	Intel	Windows
506	Asus	Notebook	8	2.00	0	0	141.211998	Intel Core i7	256	0	Intel	Windows
540	Dell	Ultrabook	8	1.20	0	1	202.372769	Intel Core i7	256	0	Intel	Windows
1222	HP	Notebook	6	2.10	0	0	141.211998	AMD Processor	0	0	AMD	Windows

1081 rows × 12 columns

Machine Learning Modeling for Laptop Price Prediction

Now we have prepared our data and hold a better understanding of the dataset. so let's get started with Machine learning modeling and find the best algorithm with the best hyperparameters to achieve maximum accuracy.

Import Libraries

```
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder'
from sklearn.metrics import r2_score,mean_absolute_error

!pip install xgboost==1.7.5

from sklearn.linear_model import LinearRegression,Ridge,Lasso
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import
RandomForestRegressor,GradientBoostingRegressor,AdaBoostRegressor,ExtraTreesRegressor
from sklearn.svm import SVRfrom xgboost import XGBRegressor
from sklearn.cluster import AffinityPropagation
```

Implement Pipeline for training and testing

Now we will implement a pipeline to streamline the training and testing process. first, we use a column transformer to encode categorical variables which is step one. After that, we create an object of our algorithm and pass both steps to the pipeline. using pipeline objects we predict the score on new data and display the accuracy.

RANDOM FOREST

Implement Pipeline for training and testing

```
step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0, 1, 7, 10, 11])
], remainder='passthrough')

# Add an imputation step to handle NaNs before the KNeighborsRegressor
step2 = SimpleImputer(strategy='mean') # Replace NaNs with the mean of the column

step3 = RandomForestRegressor(n_estimators=100,
                              random_state=3,
                              max_samples=0.5,
                              max_features=0.75,
                              max_depth=15)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2),
    ('step3', step3)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

```
R2 score 0.8765144459399266
MAE 0.1641282888254933
```

LINEAR REGRESSION

```
# ... (Your existing code for train_test_split and X_train)

step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0,1,7,10,11])
], remainder='passthrough')

# Add an imputation step to handle NaNs
step2 = SimpleImputer(strategy='mean') # Replace NaNs with the mean of the column

step3 = LinearRegression()

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2), # Impute missing values before fitting the model
    ('step3', step3)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

```
R2 score 0.8276505512108473
MAE 0.19421244887656317
```

DECISION TREE

```
# ... (Your existing code for train_test_split and X_train)

step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0,1,7,10,11])
], remainder='passthrough')

# Add an imputation step to handle NaNs before the KNeighborsRegressor
step2 = SimpleImputer(strategy='mean') # Replace NaNs with the mean of the column

step3 = DecisionTreeRegressor(max_depth=8)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2), # Impute missing values
    ('step3', step3)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

```
R2 score 0.8470529703979255
MAE 0.1747368226634528
```

STACKING

What is stacking in machine learning?

Stacking is one of the most popular ensemble machine learning techniques.

used to predict multiple nodes to build a new model and improve models performance.

```
import pandas as pd
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.metrics import r2_score, mean_absolute_error

# Add an imputation step to your pipeline
step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first', handle_unknown='ignore'), [0, 1, 7, 10, 11])
], remainder='passthrough')

step2 = SimpleImputer(strategy='mean') # Impute missing values with the mean

step3 = RandomForestRegressor(n_estimators=100,
                              random_state=3,
                              max_samples=0.5,
                              max_features=0.75,
                              max_depth=15)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2), # Add imputation step
    ('step3', step3)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))

R2 score 0.8765144459399266
MAE 0.1641282888254933
```

CONCLUSION

The development of a robust machine learning model for predicting laptop *prices has yielded valuable insights into the dynamic landscape of the tech market for SmartTech Co.* Through meticulous data exploration, preprocessing, and model development, we've identified key features influencing [(Company) (TypeName) (Ram) (Weight) (IPS) (Touchscreen) (ppi) (Cpu brand) (SSD) (HDD) (Gpu brand) (os)]

while demonstrating the model's efficacy in accurately predicting prices across various brands and specifications. Leveraging real-time prediction capabilities, SmartTech Co. can strategically position its laptops in the market, optimize pricing strategies, and swiftly respond to market shifts, ensuring competitiveness and customer satisfaction in an ever-evolving industry.

Questions to Explore:

1). Which features have the most significant impact on laptop prices?

-> The most significant features on price is 'Ram', 'Gpu brand', 'os' and 'PPI' due to it have strong correlation with price.

2). Does the brand of the laptop significantly influence its price?

-> Only applies to apple and gaming laptop. But for remaining it depends on the configuration.

3). How well does the model perform on laptops with high-end specifications compared to budget laptops?

-> Model doesn't differentiate and performs reasonably well on both ends of specifications.

4). What are the limitations and challenges in predicting laptop prices accurately?

certain brands, configurations etc.

-> Major Limitation is lack of sufficient data of most laptop companies which will adversely effect in predicting certain brands, configurations etc

Create GUI of Laptop Price Prediction Mode

!pip install gradio==4.29.0

```
[331] !pip install gradio==4.29.0

Requirement already satisfied: httpx>=0.24.1 in c:\users\l6reh\appdata\roaming\python\python311\site-packages (from gradio==4.29.0) (0.27.2)
Requirement already satisfied: huggingface-hub>=0.19.3 in c:\users\l6reh\appdata\roaming\python\python311\site-packages (from gradio==4.29.0) (0.24.6)
Requirement already satisfied: importlib-resources<7.0,>=1.3 in c:\users\l6reh\appdata\roaming\python\python311\site-packages (from gradio==4.29.0) (6.4.4)
Requirement already satisfied: Jinja2<4.0 in c:\programdata\anaconda3\lib\site-packages (from gradio==4.29.0) (3.1.3)
Requirement already satisfied: MarkupSafe==2.0 in c:\programdata\anaconda3\lib\site-packages (from gradio==4.29.0) (2.1.3)
Requirement already satisfied: matplotlib==3.0 in c:\programdata\anaconda3\lib\site-packages (from gradio==4.29.0) (3.8.0)
Requirement already satisfied: numpy==1.0 in c:\programdata\anaconda3\lib\site-packages (from gradio==4.29.0) (1.26.4)
Requirement already satisfied: orjson==3.0 in c:\users\l6reh\appdata\roaming\python\python311\site-packages (from gradio==4.29.0) (3.10.7)
Requirement already satisfied: packaging in c:\programdata\anaconda3\lib\site-packages (from gradio==4.29.0) (23.1)
Requirement already satisfied: pandas<3.0,>=1.0 in c:\programdata\anaconda3\lib\site-packages (from gradio==4.29.0) (2.1.4)
Requirement already satisfied: pillow<11.0,>=8.0 in c:\programdata\anaconda3\lib\site-packages (from gradio==4.29.0) (10.2.0)
Requirement already satisfied: pydantic>=2.0 in c:\users\l6reh\appdata\roaming\python\python311\site-packages (from gradio==4.29.0) (2.9.0)
Requirement already satisfied: pydub in c:\users\l6reh\appdata\roaming\python\python311\site-packages (from gradio==4.29.0) (0.25.1)
Requirement already satisfied: python-multipart==0.0.9 in c:\users\l6reh\appdata\roaming\python\python311\site-packages (from gradio==4.29.0) (0.0.9)
Requirement already satisfied: pyyaml<7.0,>=5.0 in c:\programdata\anaconda3\lib\site-packages (from gradio==4.29.0) (6.0.1)
Requirement already satisfied: ruff>=8.2.2 in c:\users\l6reh\appdata\roaming\python\python311\site-packages (from gradio==4.29.0) (8.6.4)
```

```
import gradio as gr
import pandas as pd
import pickle

df = pd.read_csv('laptop.csv')
pipe = pickle.load(open('pipe.pkl','rb'))
```

```
import gradio as gr
import pandas as pd
import pickle

df = pd.read_csv('laptop.csv')
pipe = pickle.load(open('pipe.pkl','rb'))

def predict_price(Company,Typellame,Ram,Weight,Touchscreen,IPS,ppi,Cpu_brand,HDD,SSD,Gpu_brand,os):
    input_df = pd.DataFrame([[Company,Typellame,Ram,Weight,Touchscreen,IPS,ppi,Cpu_brand,HDD,SSD,Gpu_brand,os]],columns=['Company','Typellame','Ram','Weight',

# Preprocess the input
    Company = input_df['Company'].map({'Apple':0,'Razer':1,'Mediacom':2,'Samsung':3,'Toshiba':4,'MSI':5,'Vero':6,'Dell':7,'Huawei':8,'Chuvi':9,'Fujitsu':10})
    input_df['Company'] = Company

    Typellame = input_df['Typellame'].map({'Ultrabook':0,'Notebook':1,'Netbook':2,'Gaming':3,'2 in 1 Convertible':4,'Workstation':5})
    input_df['Typellame'] = Typellame

    Cpu_brand = input_df['Cpu_brand'].map({'Intel Core i7':0,'Intel Core i5':1,'Intel Core i3':2,'Other Intel Processor':3,'AMD Processor':4})
    input_df['Cpu_brand'] = Cpu_brand

    Gpu_brand = input_df['Gpu_brand'].map({'Intel':0,'Nvidia':1,'AMD':2,'ARM':3})
    input_df['Gpu_brand'] = Gpu_brand

    os = input_df['os'].map({'Mac':0,'Others/No OS/Linux':1,'Windows':2})
    input_df['os'] = os

    prediction = np.exp(pipe.predict(input_df))
    return prediction[0]
```

Define the input components for the interface

```
# Define the input components for the interface
new_var = [
    gr.Dropdown(choices=list(df['Company'].unique()), label="Company"),
    gr.Dropdown(choices=list(df['TypeName'].unique()), label="TypeName"),
    gr.Slider(4, 64, step=4, label="Ram"),
    gr.Slider(0.5, 5.0, step=0.1, label="Weight"),
    gr.Radio(choices=[0, 1], label="Touchscreen"),
    gr.Radio(choices=[0, 1], label="IPS"),
    gr.Slider(50, 500, step=1, label="ppi"),
    gr.Slider(0, 2000, step=1, label="HDD"),
    gr.Slider(0, 2000, step=1, label="SSD"),
    gr.Dropdown(choices=list(df['OpSys'].unique()), label="os") # Changed 'os' to 'OpSys'
]
inputs = new_var

# Define the output component for the interface
outputs = gr.Number(label="Price")

# Create the Gradio interface
iface = gr.Interface(fn=predict_price, inputs=inputs, outputs=outputs)

# Launch the interface
iface.launch()
share=True
```

Laptop Prediction GUI

Running on local URL: <http://127.0.0.1:7861>

To create a public link, set ``share=True`` in ``launch()``.

Company

TypeName

Ram

4

Weight

0.5

Touchscreen

☐ 0

☐ 1

IPS

☐ 0

☐ 1

Price

0

Flag

THANK YOU

REHAN ASLAM KHAN 

+91 8968832194 

16rehan687@gmail.com 

<https://github.com/REHAN-S8882> Github.com 